

Principal Component Analysis with Python Programming


Chrysovalantou Kalaitzidou

March 2017

University of Crete Faculty of Medicine
Bioinformatics M.Sc.

Methods in Bioinformatics

First Assignment



Contents

1	Introduction	iii
2	Theoretical Background	1
2.1	Principal Component Analysis	1
2.1.1	Maximun Variance Formulation	1
2.2	Probabilistic PCA	3
2.3	Kernel PCA	6
3	Analysis	8
3.1	Theoretical Problem	8
3.1.1	Part (i): Synthetic Data	8
3.1.2	Part (ii): Plot of Eigenvalue spectrum	18
3.2	Practical Problem	19
4	Python Codes	23
5	Suggested Bibliography	31

1 Introduction

The current report is part of the first assignment in class *Methods in Bioinformatics* and refers to ***Principal Component Analysis***, with respect to its methods and algorithms. Certain data sets have been provided for analysis with each method implemented in Python programming language. In the following Chapters the reader can find some Theoretical Notes for the methods and their algorithms for the implementation, the Results derived from data sets' analysis along with some discussion and, at last, the Python Code.

The data set that has been given for the *Practical Part* of analysis, acquired from *NCBI* database, refers to organism *Mus musculus* and a particular study of livers of *C57BL/6J mice* fed a high fat diet for up to 24 weeks. Significant body weight gain was observed after 4 weeks. Their results provide insight into the effect of high fat diets on metabolism in the liver. For more information, please visit <https://www.ncbi.nlm.nih.gov/sites/GDSbrowser?acc=GDS6248>

Chrysovalantou Kalaitzidou

2 Theoretical Background

2.1 Principal Component Analysis

Principal component analysis, (PCA) is probably the oldest and best known of the techniques of *Multivariate Analysis*. It was first introduced by Pearson (1901), and developed independently by Hotelling (1933). Like many multivariate methods, it was not widely used until the advent of electronic computers, but it is now well entrenched in virtually every statistical computer package.

The central idea of principal component analysis is to reduce the dimensionality of a data set in which there are a large number of interrelated variables, while retaining as much as possible of the variation present in the data set. This reduction is achieved by transforming to a new set of variables, the *Principal Components*, which are uncorrelated, and which are ordered so that the first few retain most of the variation present in all of the original variables. Computation of the Principal Components reduces to the solution of an eigenvalue-eigenvector problem for a positive-semidefinite symmetric matrix. Thus, the definition and computation of Principal Components are straightforward but, as will be seen, this apparently simple technique has a wide variety of different applications, as well as a number of different derivations.

There are two commonly used definitions of *PCA* that give rise to the same algorithm. *PCA* can be defined as the orthogonal projection of the data onto a lower dimensional linear space, known as the *Principal Subspace*, such that the variance of the projected data is *maximized* (Hotelling, 1933). Equivalently, it can be defined as the linear projection that *minimizes* the average projection cost, defined as the mean squared distance between the data points and their projections. In the current project we discuss the *Maximum variance formulation* only.

At last, we should cite that *PCA* is sometimes used as a preliminary to, or in conjunction with, other statistical techniques, especially with three well-known multivariate techniques, namely *Discriminant Analysis*, *Cluster Analysis* and *Canonical Correlation Analysis*.

2.1.1 Maximun Variance Formulation

Consider a data set of observations x_n where $n = 1, \dots, N$, and x_n is an Euclidean variable with dimensionality D . Our goal is to project the data onto a space having dimensionality $M < D$ while maximizing the variance of the projected data. To begin with, consider the projection onto a one-dimensional space ($M = 1$). We can define the direction of this space using a D -dimensional vector u_1 , which for convenience (and without loss of generality) we shall choose to be a unit vector so that

$$u_1^T u_1 = 1$$

(note that we are only interested in the direction defined by u_1 , not in the magnitude of u_1 itself). Each data point x_n is then projected onto a scalar value

$$u_1^T x_n$$

The mean of the projected data is $u_1 \bar{x}$ where \bar{x} is the sample set mean given by

$$\bar{x} = \sum_{n=1}^N \frac{x_n}{N}$$

and the variance of the projected data is given by

$$\frac{1}{N} \sum_{n=1}^N (u_1^T x_n - u_1^T \bar{x})^2 = u_1^T S u_1$$

where S is the data covariance matrix defined by

$$S = \frac{1}{N} \sum_{n=1}^N (x_n - \bar{x})(x_n - \bar{x})^T$$

We now maximize the projected variance $u_1^T S u_1$ with respect to u_1 . Clearly, this has to be a constrained maximization to prevent $\|u_1\| \rightarrow \infty$. The appropriate constraint comes from the normalization condition $u_1^T u_1 = 1$. To enforce this constraint, we introduce a *Lagrange multiplier* that we shall denote by λ_1 , and then make an unconstrained maximization of

$$u_1^T S u_1 + \lambda_1(1 - u_1^T u_1)$$

By setting the derivative with respect to u_1 equal to zero, we see that this quantity will have a stationary point when

$$S u_1 = \lambda_1 u_1$$

which says that u_1 must be an *eigenvector* of S . If we left-multiply by u_1^T and make use of $u_1^T u_1 = 1$, we see that the variance is given by

$$u_1^T S u_1 = \lambda_1$$

and so the variance will be a *maximum when we set u_1 equal to the eigenvector having the largest eigenvalue λ_1* . This eigenvector is known as the ***First Principal Component***.

We can define additional Principal Components in an incremental fashion by choosing each new direction to be that which maximizes the projected variance amongst all possible directions orthogonal to those already considered. If we consider the general case of an M -dimensional projection space, the optimal linear projection for which the variance of the projected data is maximized is now defined by the M eigenvectors u_1, \dots, u_M of the data *Covariance matrix* S corresponding to the M largest eigenvalues $\lambda_1, \dots, \lambda_M$. This is easily shown using proof by induction.

To summarize, Principal Component Analysis involves:

- Evaluating the mean \bar{x} of the data set.
- Evaluating the Covariance matrix S of the data set.
- Finding the M eigenvectors of S corresponding to the M largest eigenvalues.

Note that the computational cost of computing the full eigenvector decomposition for a matrix of size $D \times D$ is $O(D^3)$. If we plan to project our data onto the first M Principal Components, then we only need to find the first M eigenvalues and eigenvectors. This can be done with more efficient techniques or alternatively we can make use of the ***EM algorithm***.

2.2 Probabilistic PCA

The formulation of *PCA* discussed in the previous section was based on a linear projection of the data onto a subspace of lower dimensionality than the original data space. We now show that PCA can also be expressed as the **Maximum Likelihood** solution of a *Probabilistic* latent variable model. This reformulation of PCA, known as **Probabilistic PCA**, brings several advantages compared with conventional PCA:

- Probabilistic PCA represents a constrained form of the Gaussian distribution in which the number of free parameters can be restricted while still allowing the model to capture the dominant correlations in a data set.
- We can derive an *EM Algorithm* for PCA that is computationally efficient in situations where only a few leading eigenvectors are required and that avoids having to evaluate the data covariance matrix as an intermediate step.
- The combination of a probabilistic model and EM allows us to deal with missing values in the data set.
- Mixtures of probabilistic PCA models can be formulated in a principled way and trained using the EM Algorithm.

Probabilistic PCA is a simple example of the linear-Gaussian framework, in which all of the marginal and conditional distributions are *Gaussian*. We can formulate Probabilistic PCA by first introducing an explicit latent variable z corresponding to the Principal Component Subspace. Next we define a Gaussian prior distribution $p(z)$ over the latent variable, together with a Gaussian conditional distribution $p(x|z)$ for the observed variable x conditioned on the value of the latent variable. Specifically, the prior distribution over z is given by a zero-mean unit-covariance Gaussian

$$p(z) = \mathcal{N}(z|0, I)$$

Similarly, the Conditional Distribution of the observed variable x , conditioned on the value of the latent variable z , is again Gaussian, of the form

$$p(x|z) = \mathcal{N}(x|Wz + \mu, \sigma^2 I)$$

in which the mean of x is a general linear function of z governed by the $D \times M$ matrix W and the D dimensional vector μ . Note that this factorizes with respect to the elements of x , in other words this is an example of the naive *Bayes model*. **The columns of W span a linear subspace within the data space that corresponds to the principal subspace.**

The other parameter in this model is the scalar σ^2 governing the *Variance* of the Conditional Distribution. Note that there is no loss of generality in assuming a zero mean, unit Covariance Gaussian for the latent Distribution $p(z)$ because a more general Gaussian Distribution would give rise to an equivalent Probabilistic Model. We can view the Probabilistic PCA model from a generative viewpoint in which a sampled value of the observed variable is obtained by first choosing a value for the latent variable and then sampling the observed variable conditioned on this latent value. Specifically, the D dimensional observed variable x is defined by a linear transformation of the M dimensional latent variable z plus additive Gaussian '*noise*', so that

$$x = Wz + \mu + \epsilon$$

where z is an M dimensional Gaussian latent variable, and ϵ is an D dimensional zero-mean Gaussian distributed *noise* variable with Covariance $\sigma^2 I$:

$$z \sim \mathcal{N}(0, I)$$

and

$$\epsilon \sim \mathcal{N}(0, \sigma^2 I)$$

The introduction of a probabilistic model facilitates the use of the **Expectation Maximization (EM) Algorithm** (Dempster et al. 1977) to estimate the latent variables. We first present the *PPCA-EM* Algorithm for complete data set. This framework was established by Tipping Bishop (1999) using an iterated, two-step process.

- **The Expectation (E) step:**

the hidden (unknown) variables are estimated from the observations and the current values of the parameters. Statistical moments of the latent variables, z_n and $z_n z_n^T$, are estimated using $p(z_n|x_n, W, \sigma^2)$, the conditional probability density of z_n given the observations x_n and the current values of W and σ^2 , (Little and Rubin 1987). Following the probability assumptions for z and ϵ , one obtains,

$$E[z_n] = (W^T W + \sigma^2 I)^{-1} W^T (x_n - \bar{x})$$

and

$$E[z_n z_n^T] = \sigma^2 M^{-1} + E[z_n] E[z_n]^T$$

where M is the $M \times M$ matrix

$$M = W^T W + \sigma^2 I$$

- **The Maximization (M) step:**

new estimates of the parameters, W and σ^2 , are computed by maximizing the conditional expectation of the log-likelihood, ℓ , with respect to the conditional probability density of the unknown variables z_n given the known variables x_n , $p(z_n|x_n, W, \sigma^2)$. The log-likelihood is defined in terms of the joint probability of the observed variables and the latent variables:

$$\ell = \sum_{n=1}^N \ln p(x_n, z_n)$$

So, we obtain the M - step equations

$$W_{new} = \left(\sum_{n=1}^N (x_n - \bar{x}) E[z_n]^T \right) \left(\sum_{n=1}^N E[z_n z_n^T] \right)^{-1}$$

and

$$\sigma_{new}^2 = \frac{1}{ND} \sum_{n=1}^N (\|x_n - \bar{x}\|^2 - 2E[z_n]^T W_{new}^T (x_n - \bar{x}) + Tr(E[z_n z_n^T] W_{new}^T W_{new}))$$

The *EM* Algorithm for Probabilistic PCA proceeds by initializing the parameters and then alternately computing the sufficient statistics of the latent space posterior distribution using the *E-step* equations and revising the parameter values using the equations in the *M-step*.

Please consider that the calculations from which these equations arise are beyond the purpose of this report, so they are not further discussed. The reader can find details in references and the suggested textbooks.

Another elegant feature of the EM approach is that we can take the limit $\sigma^2 \rightarrow 0$, corresponding to standard PCA, and still obtain a valid EM-like algorithm (Roweis,1998). From E and M steps above, we see that the only quantity we need to compute in the E step is $E[z_n]$. Furthermore, the M step is simplified because $M = W^T W$. To emphasize the simplicity of the algorithm, let us define \tilde{X} to be a matrix of size $D \times N$ whose n^{th} row is given by the vector $x_n - \bar{x}$ and similarly define Ω to be a matrix of size $D \times M$ whose n^{th} row is given by the vector $E[z_n]$. The **E step** of the EM Algorithm for PCA then becomes

$$\Omega = (W_{old}^T W_{old})^{-1} W_{old}^T \tilde{X}$$

and the **M step** takes the form

$$W_{new} = \tilde{X}^T \Omega^T (\Omega \Omega^T)^{-1}$$

These equations have a simple interpretation as follows. The E step involves an orthogonal projection of the data points onto the current estimate for the principal subspace. Correspondingly, the M step represents a re-estimation of the principal subspace to minimize the squared reconstruction error in which the projections are fixed.

In both cases, $\sigma^2 \rightarrow 0$ or not, the E and M steps are repeated until a suitable *convergence criterion* is satisfied.

2.3 Kernel PCA

Many linear parametric models can be re-cast into an equivalent ‘dual representation’ in which the predictions are also based on linear combinations of a *Kernel Function* evaluated at the training data points. For models which are based on a fixed *nonlinear feature space* mapping $\phi(x)$, the **Kernel Function** is given by the relation

$$k(x_i, x_j) = \phi(x_i)^T \phi(x_j)$$

From this definition we see that the Kernel is a symmetric function of its arguments so that $k(x_i, x_j) = k(x_j, x_i)$. Given any algorithm that can be expressed solely in terms of dot products, this *trick* allows us to construct different nonlinear versions of it

The concept of a Kernel formulated as an inner product in a feature space allows us to build interesting extensions of many well-known algorithms by making use of the *Kernel Trick*, also known as *Kernel substitution*. The general idea is that, if we have an algorithm formulated in such a way that the input vector x enters only in the form of scalar products, then we can replace that scalar product with some other choice of Kernel.

Popular Kernels:

- Gaussian Kernel:

$$K(x_i, x_j) = \exp(-\gamma \|x_i - x_j\|^2)$$

- Polynomial Kernel:

$$K(x_i, x_j) = (1 + x_i x_j)^p$$

- Hyperbolic Tangent:

$$K(x_i, x_j) = \tanh(x_i x_j + \delta)$$

When this technique of kernel substitution is applied to principal component analysis, we obtain a nonlinear generalization called *Kernel PCA* (Schölkopf et al., 1998).

Consider a data set x_n of observations, where $n = 1, \dots, N$, in a space of dimensionality D . In order to keep the notation uncluttered, we shall assume that we have already subtracted the sample mean from each of the vectors x_n , so that $\sum_{n=1}^N x_n = 0$. The first step is to express conventional PCA in such a form that the data vectors x_n appear only in the form of the scalar products $x_n^T x_m$. Recall that the principal components are defined by the eigenvectors u_i of the covariance matrix $S u_i = \lambda_i u_i$ where $i = 1, \dots, D$. Here $D \times D$ Covariance matrix S is defined by

$$S = \frac{1}{N} \sum_{n=1}^N x_n x_n^T$$

and the eigenvectors are normalized such that $u_i^T u_i = 1$.

Now consider a nonlinear transformation $\phi(x)$ into an M -dimensional feature space, so that each data point x_n is thereby projected onto a point $\phi(x_n)$. We can now perform standard PCA in the feature space, which implicitly defines a *nonlinear Principal Component* model in the original data space. The $M \times M$ sample covariance matrix in feature space is given by

$$C = \frac{1}{N} \sum_{n=1}^N \phi(x_n) \phi(x_n)^T$$

and its eigenvector expansion is defined by $C v_i = \lambda_i v_i$, $i = 1, \dots, M$.

Our goal is to solve this eigenvalue problem without having to work explicitly in the feature space.

From the definition of C , the eigenvector equations tells us that v_i satisfies

$$\frac{1}{N} \sum_{n=1}^N \phi(x_n)(\phi(x_n^T)v_i) = \lambda_i v_i$$

and so we see that (provided $\lambda_i > 0$) the vector v_i is given by a linear combination of the $\phi(x_n)$ and so can be written in the form

$$v_i = \sum_{n=1}^N \alpha_{in} \phi(x_n)$$

Substituting this expansion back into the eigenvector equation, we obtain

$$\frac{1}{N} \sum_{n=1}^N \phi(x_n) \phi(x_n)^T \sum_{m=1}^N \alpha_{im} \phi(x_m) = \lambda_i \sum_{n=1}^N \alpha_{in} \phi(x_n)$$

The key step is now to express everything in terms of the kernel function $k(x_n, x_m) = \phi(x_n^T) \phi(x_m)$, which we do by multiplying both sides by $\phi(x_l)^T$ to give

$$\frac{1}{N} \sum_{n=1}^N k(x_l, x_n) \sum_{m=1}^M \alpha_{im} k(x_n, x_m) = \lambda_i \sum_{n=1}^N \alpha_{in} k(x_l, x_n)$$

This can be written in matrix notation as $K^2 a_i = \lambda_i N K a_i$. where a_i is an N -dimensional column vector with elements a_{ni} for $n = 1, \dots, N$. We can find solutions for a_i by solving the following eigenvalue problem: $K a_i = \lambda_i N a_i$. The normalization condition for the coefficients a_i is obtained by requiring that the eigenvectors in feature space be normalized ($1 = v_i^T v_i$). Having solved the eigenvector problem, the resulting principal component projections can then also be cast in terms of the kernel function so that, the projection of a point x onto eigenvector i is given by:

$$y_i(x) = \phi(x)^T v_i = \sum_{n=1}^N \alpha_{in} \phi(x)^T \phi(x_n) = \sum_{n=1}^N \alpha_{in} k(x, x_n)$$

and so again is expressed in terms of the kernel function. In the original D -dimensional x space there are D orthogonal eigenvectors and hence we can find at most D linear principal components. The dimensionality M of the feature space, however, can be much larger than D (even infinite), and thus we can find a number of nonlinear principal components that can exceed D . Note, however, that the number of nonzero eigenvalues cannot exceed the number N of data points, because (even if $M > N$) the covariance matrix in feature space has rank at most equal to N . This is reflected in the fact that kernel PCA involves the eigenvector expansion of the $N \times N$ matrix K . So far we have assumed that the projected data set given by $\phi(x_n)$ has zero mean, which in general will not be the case. We cannot simply compute and then subtract off the mean, since we wish to avoid working directly in feature space, and so again, we formulate the algorithm purely in terms of the kernel function. The projected data points after centralizing, denoted $\bar{\phi}(x_n)$ are given by

$$\bar{\phi}(x_n) = \phi(x_n) - \frac{1}{N} \sum_{l=1}^N \phi(x_l)$$

and the corresponding elements of the Gram matrix, expressed in matrix notation, are given by:

$$\bar{K} = K - 1_N K - K 1_N + 1_N K 1_N$$

where 1_N denotes the $N \times N$ matrix in which every element takes the value $\frac{1}{N}$. Thus we can evaluate \bar{K} using only the Kernel function and then use \bar{K} to determine the eigenvalues and eigenvectors.

One obvious disadvantage of kernel PCA is that it involves finding the eigenvectors of the $N \times N$ matrix \bar{K} rather than the $D \times D$ matrix S of conventional linear PCA, and so in practice for large data sets approximations are often used. Finally, we note that in standard linear PCA, we often retain some reduced number $L < D$ of eigenvectors and then approximate a data vector x_n by its projection onto the L -dimensional principal subspace, defined by

$$\hat{x} = \sum_{i=1}^L (x_n^T u_i) u_i$$

3 Analysis

3.1 Theoretical Problem

3.1.1 Part (i): Synthetic Data

```
1 from sklearn.datasets import make_circles
2 X, y = make_circles(n_samples=1000, factor=0.3, noise=0.05)
```

We consider here the nonlinear problem *concentric circles*. We assume two class problem where the triangle shapes represent one class and the circle shapes represent another class, respectively:

We shall perform PCA and PCa Kernel Methods and compare the results.

PCA method, either with *Eigendecomposition* or *SVD* Algorithms, and **Probabilistic PCA**, either with $\sigma^2 \neq 0$ or $\sigma^2 \rightarrow 0$: None of these techniques is able to produce results suitable for training a linear classifier.

Initial plot of the *concentric circles*

```
1 from sklearn.datasets import make_circles
2 from matplotlib import pyplot as plt
3
4 X, y = make_circles(n_samples=1000, factor=0.3, noise=0.05)
5
6 plt.scatter(X[y==0,0],X[y==0,1],color='red',marker='^',alpha=0.5,label='Circle_01')
7 plt.scatter(X[y==1,0],X[y==1,1],color='blue',marker='o',alpha=0.5,label='Circle_02')
8 plt.grid(True)
9
10 plt.legend(numpoints =1,loc='lower right')
11 plt.title('Concentric circles')
12 plt.savefig("Initial plot")
13 plt.show()
```

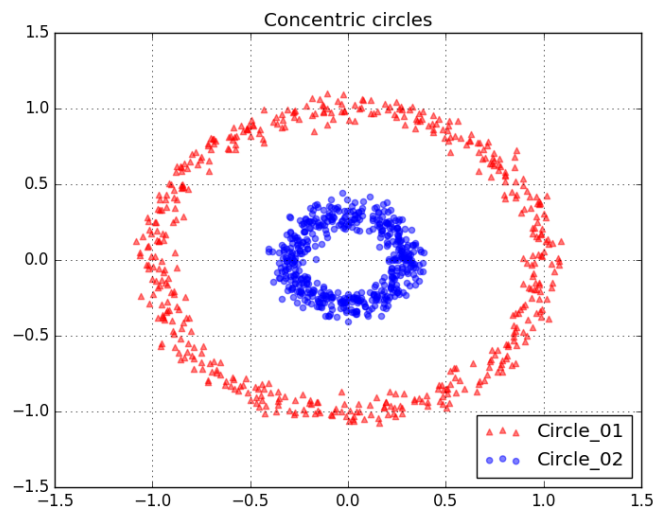


Figure 1: Concentric Circles Initial Plot

PCA with Eigendecomposition




Figure 2: (a) Projection in 1d

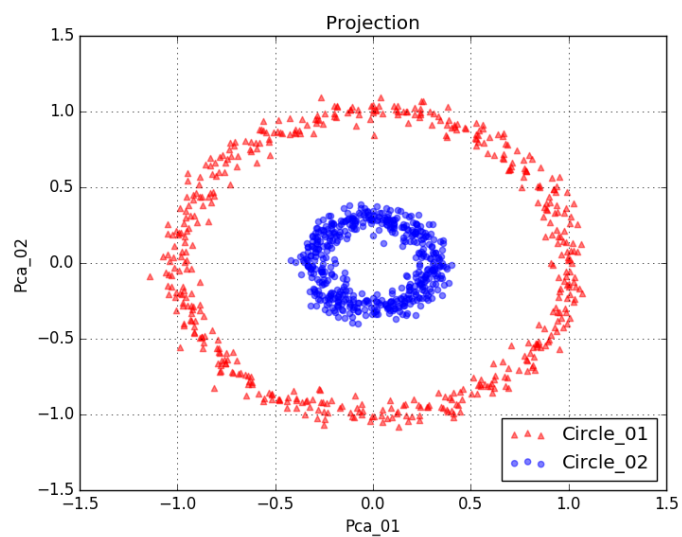


Figure 3: (b) Projection in 2d

PCA with SVD

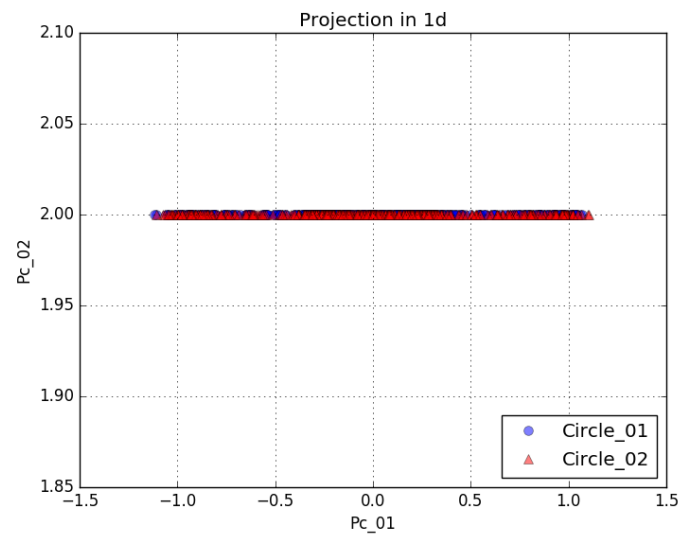


Figure 4: (a) Projection in 1d




Figure 5: (b) Projection in 2d

PPCA with $\sigma^2 = 0$




Figure 6: (a) Projection in 1d




Figure 7: (b) Projection in 2d

PPCA with $\sigma^2 \neq 0$




Figure 8: (a) Projection in 1d




Figure 9: (b) Projection in 2d

Kernel methods

As we already discussed, there is a variety of different Kernel Functions. Choosing the most appropriate kernel highly depends on the problem at hand – and fine tuning its parameters can easily become a tedious and cumbersome task. The choice of a Kernel depends on the problem at hand because it depends on what we are trying to model. A polynomial kernel, for example, allows us to model feature conjunctions up to the order of the polynomial. On the other hand, Radial basis functions, such as Gaussian Kernel, allow us to pick out circles (or hyperspheres). The motivation behind the choice of a particular kernel can be very intuitive and straightforward depending on what kind of information we are expecting to extract about the data.

Here, we perform Kernel PCA for the three Kernels *Gaussian*, *Polynomial* and *Hyperbolic Tangent*.

Gaussian Kernel PCA, as it was expected, projected the data onto a new subspace where the two classes become linearly separable:

Kernel PCA with Gaussian Kernel

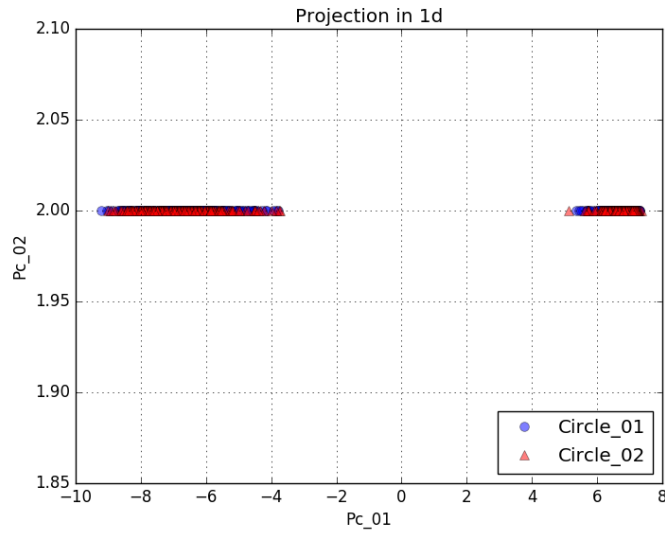


Figure 10: (a) Projection in 1d $\gamma = 5$

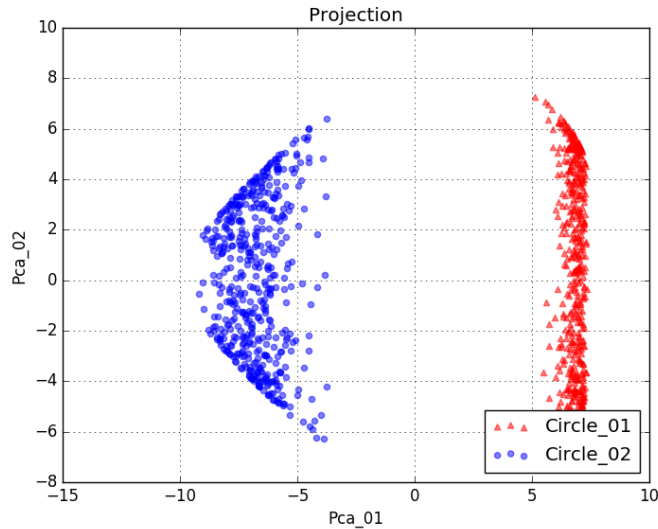


Figure 11: (b) Projection in 2d $\gamma = 5$




Figure 12: (c) Projection in 3d $\gamma = 5$

However, the other two Kernel functions *Polynomial* and *Hyperbolic Tangent* do not behave in the same way:

Kernel PCA with Polynomial Kernel




Figure 13: Projection in 3d and $p = 4$

Kernel PCA with Tangent Kernel

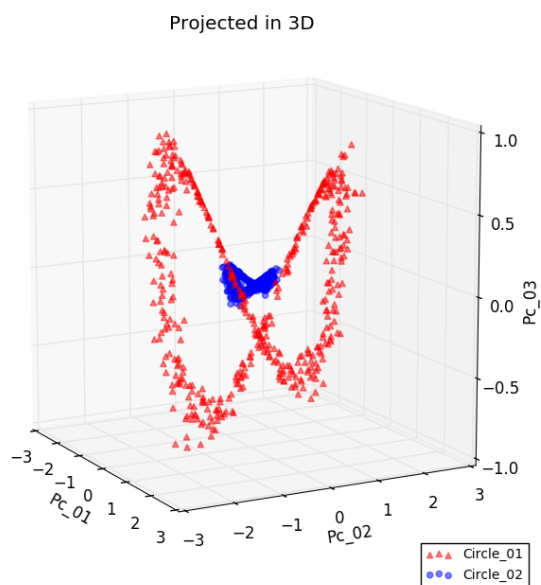


Figure 14: Projection in 3d and $\delta = 2$

Now, we change the *noise* parameter in the concentric circles to *noise* = 3 and perform *Gaussian Kernel* :

Kernel PCA with Gaussian Kernel

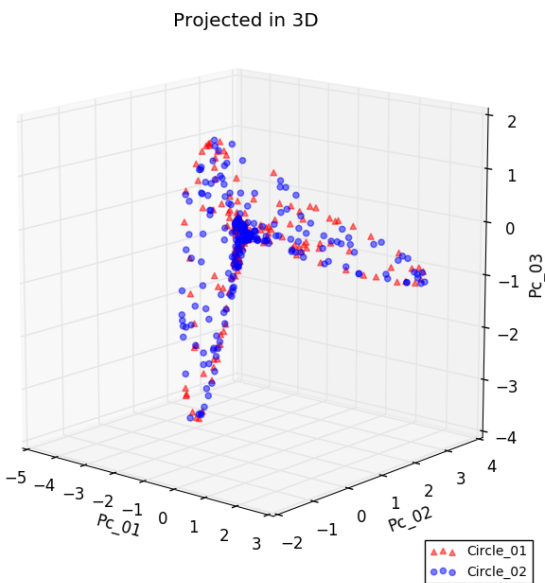


Figure 15: (a) Projection in 3d $\gamma = 5$

For $\gamma = 20, 40, 60$ we have respectively:

Kernel PCA with Gaussian Kernel




Figure 16: Projection in 3d $\gamma = 20$

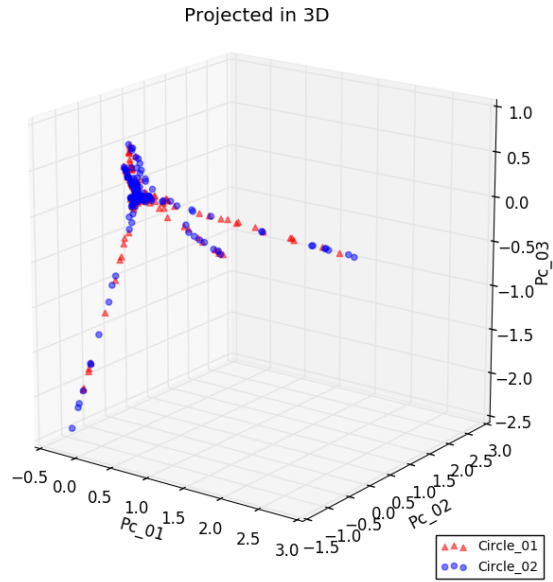


Figure 17: Projection in 3d $\gamma = 40$

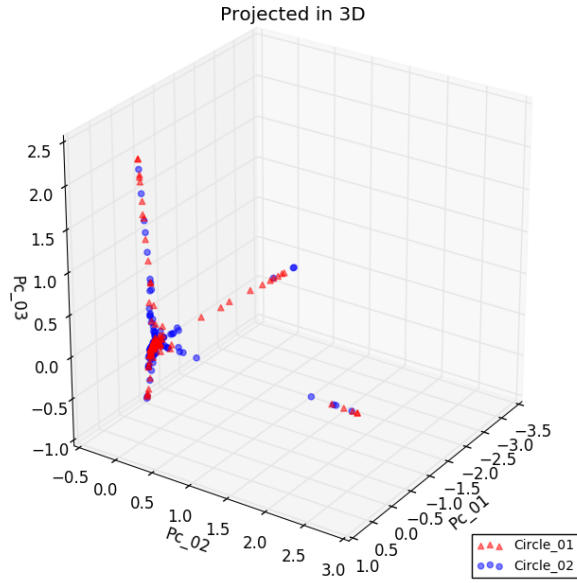


Figure 18: Projection in 3d $\gamma = 60$

Increasing the *noise* in the samples data, increases the complexity between them ,so it is difficult to make them separate. It is interesting, yet, to see that for different values of γ the *projection* tends to be orthogonal with respect to each principal component.

3.1.2 Part (ii): Plot of Eigenvalue spectrum

In this part, we generate a dataset X which is randomly Gaussian distributed with mean a 10×1 vector with 1 in every position and 10 samples and make a plot of the eigenvalues in descending order. Then, we keep only 5 samples and inspect the eigenvalues. Below is a combined plot for both of the eigenvalues spectrum: While reducing dimensionality $10 \rightarrow 5$ *columns*, the number of *observations* is larger than the number of *samples*, in other words $D < N$. So, what we expect to see is that, in the second spectrum, $D - N$ eigenvalues tend $\rightarrow 0$ while all the others are *raised*, as a result of the *noise* in Covariance Matrix.

The reader is encouraged to read more about the *noise* in covariance matrices and the determination of the eigenvalue spectrum when the number of observations is limited in the [5].

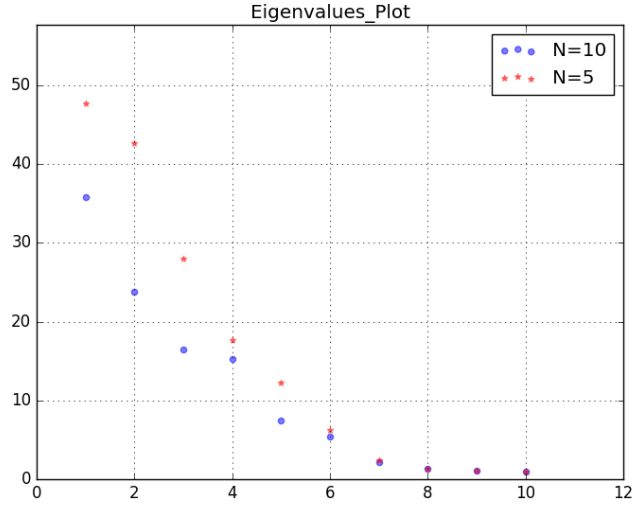


Figure 19: Eigenvalues Spectrum

Plot of the eigenvalue spectrum for the 10 samples set (circle shapes) and for the 5 samples set (star shapes).

3.2 Practical Problem

The aim of this exercise is to perform Principal Component Analysis in a real data set. The analysis of livers of *C57BL/6J* mice fed a high fat diet for up to 24 weeks has shown significant body weight gain was observed after 4 weeks. The results provide insight into the effect of high-fat diets on metabolism in the liver. We have 51 observations divided in 3 categories: *Baseline*, *Normal Diet* and *High– Fat Diet*. The aim is to see if the implemented *PCA Algorithm* is good approach for this data set especially in comparison to PCA analysis with the particular *built– in* module in *Python*.

The number of the observations, as mentioend above, is 51, while the number of dimensions 45281. We perform Principal Component Analysis, using the *EM Algorithm*, for $\sigma^2 = 0$ and $\sigma^2 > 0$. As we mentioned in theory, In both cases, $\sigma^2 \rightarrow 0$ or not , the *E* and *M* steps of the algorithm are repeated until a suitable *convergence criterion* is satisfied. Here, we make use of **Root Mean Squared Error (RMSE)** for the convergence of the *W* Matrix.

$$RMSE = \sqrt{\frac{1}{N} \sum_{i=1}^n (y_i - \hat{y}_i)^2}$$

Basically, it is the square root of the mean of the square of all of the error. The use of *RMSE* is very common and it makes an excellent general purpose error metric for numerical predictions. Compared to the similar Mean Absolute Error, RMSE amplifies and severely punishes large errors. For the convergence of σ^2 we make use of

$$\|\sigma_{new}^2 - \sigma_{old}^2\|$$

EM Algorithm with $\sigma^2 = 0$




Figure 20: Projection in 2d, and number of iterations: 10




Figure 21: Projection in 2d, and number of iterations: 10

EM Algorithm with $\sigma^2 = 0$




Figure 22: Projection in 2d, and number of iterations: 20

The previous plot is one of the various ones that gained after performing, all displaying the same projection ,only rotated a little bit. But, the key point is that it is very similar to the projection of built in function for PCA algorithm:

PCA Analysis




Figure 23: Projection in 2d using PCA built algorithm

Number of Iterations refers to the number of calls of EM steps, so as to get a *mean* W_{new} matrix.

EM Algorithm with $\sigma^2 > 0$

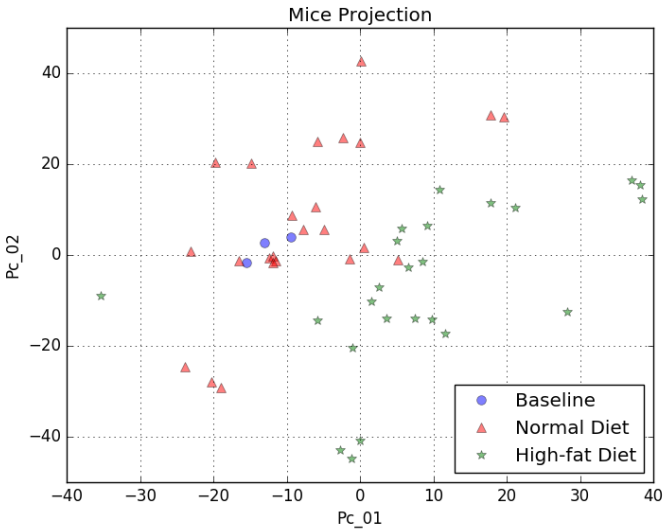


Figure 24: Projection in 2d, and number of iterations: 10

EM Algorithm with $\sigma^2 > 0$




Figure 25: Projection in 2d, and number of iterations: 20

PCA Analysis

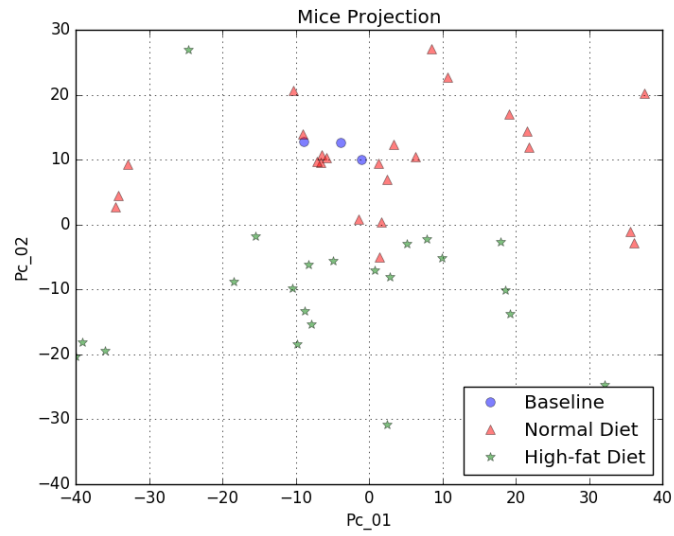


Figure 26: Projection in 2d using PCA built algorithm

In this case as well, we get the same projection close to PCA Algorithm, only the subspace is rotated.

4 Python Codes

File: main.py

```
1 import time
2 import os
3 from Principal_Methods import*
4
5 Problem = input("Choose either Theoretical or Practical Problem.\n Enter A or B for Theoretical or
6 Practical respectively:")
7
8 if Problem == 'A':
9     print("****___Theoretical Problem has been selected___****\n")
10
11     Sub_quest = input("Choose Subquestion for the Theoretical part\n Enter A for Subquestion (i) and B
12     for Subquestion (ii): ")
13
14     if Sub_quest == 'A':
15         print("****___Theoretical Part: Subquestion (i)___****")
16
17         Method = input("Choose Method: Enter PCA for Principal Component Analysis,\nEM for Probabilistic
18         PCA and \nKERNEL for Kernel Method: ")
19         Dimension = int(input("Give Dimensionality of Projection\n Notice that for either PCA or EM
20         should be 1 or 2: "))
21         M = Dimension
22
23         ## Data given for the theoretical problem . Subquestion (i) with different noises:
24
25         n = input("Enter A for 0.05 noise: ")
26         if n == 'A':
27             X, y = make_circles(n_samples=1000,noise=0.05, factor=0.3)
28         else:
29             X, y = make_circles(n_samples=1000,noise=3, factor=0.3)
30
31         Data = X.T
32
33         if Method == 'PCA':
34             PCA(Data,y,M)
35         elif Method == 'EM':
36             PPCA(Data,y,M)
37         else:
38             KERNEL(Data,y,M)
39     else:
40         print("****___Theoretical Part: Subquestion (ii)___****\n")
41
42         ## Data given for the theoretical problem . Subquestion (ii):
43
44         mean = np.array([1,1,1,1,1,1,1,1,1,1])#; print(mean.shape)
45         cov = np.eye((10))#; print(cov.shape)
46
47         Xa = np.random.multivariate_normal(mean,cov,10)
48         Xb = Xa[:,0:5]
49
50         Theoretical_II(Xa,Xb)
51
52 else:
53     y = [0]
54     print("****___Practical Problem has been selected___****\n")
55
56     print("\nData set should be downloaded automatically and the process shall begin.\n")
57
58     if not os.path.exists("Final.txt"):
59
60         Filename = os.system('wget
61         ftp://ftp.ncbi.nlm.nih.gov/geo/datasets/GDS6nnn/GDS6248/soft/GDS6248.soft.gz')
62
63         os.system('gunzip <GDS6248.soft.gz> Data_set.txt')
64         os.system('grep -i ILMN Data_set.txt > Data.txt')
65         os.system('cut -f3- Data.txt > Final.txt')
66     else:
67         print('Skipping file download, Data file exists...')
68
69     X = np.loadtxt("Final.txt") ## Constructing the Data Array
70     print(X.shape)
71
72     Method = input("Choose Method: Enter PCA for Principal Component Analysis, EM for Probabilistic PCA
73     and KERNEL for Kernel Method: ")
74     Dimension = int(input("Give Dimensionality of Projection:"))
75     M = Dimension
76
77     if Method == 'PCA':
78         PCA(X,y,M)
79     elif Method == 'EM':
80         PPCA(X,y,M)
81     else:
82         KERNEL(X,y,M)
```

```
78 |
79 | ##### PCA with Built in class
80 |
81 | from sklearn.decomposition import PCA as pca
82 |
83 | n_components = Dimension
84 | my_pca = pca(n_components)
85 |
86 | Projected_Data = my_pca.fit_transform(X.T).T
87 |
88 | if Dimension == 2:
89 |     Practical_Plots(Projected_Data)
90 | else:
91 |     D_Plots(Projected_Data)
92 |
93 | print("\n\n*****____End of Process____*****\n\n")
94 |
95 |
96 | Time = time.process_time()
97 | print(Time)
```

File: PrincipalMethods.py

```

1 ##### File including All Methods for Principal Component Analysis
2
3 import numpy as np
4 from random import random
5 from scipy import linalg
6 from matplotlib import pyplot as plt
7 from mpl_toolkits.mplot3d import Axes3D
8 from mpl_toolkits.mplot3d import proj3d
9 from sklearn.datasets import make_circles
10 from sklearn.metrics import mean_squared_error
11
12
13 ##### Principal Component Analysis #####
14
15 def PCA(Arr,y,M):
16     print("\n#####__Principal Component Analysis__#####\n")
17
18     D = Arr.shape[0]
19     N = Arr.shape[1]
20
21     ## Mean Vector:
22
23     mean_vector = np.empty([D,1])
24     for i in range(D):
25         mean_vector[i] = np.mean(Arr[i,:])
26     # print(mean_vector.shape)
27
28     mean_d = np.repeat (mean_vector,N,axis=1)      # The mean vector with DxN shape
29     X = Arr- mean_d                                # Normalize the Data matrix, X: DxN
30
31     ## Scatter Matrix: is used to estimate the Covariance matrix of a multivariate normal distribution
32
33     Scatter_Matrix = np.empty([D,D])
34     for i in range(X.shape[1]):
35         Scatter_Matrix += (X[:,i].reshape(D,1)).dot((X[:,i].reshape(D,1)).T)
36         #Scatter_Matrix += (X[:,i].reshape(D,1) - mean_vector).dot((X[:,i].reshape(D,1) -
37             mean_vector).T)
38     print('Scatter Matrix:\n', Scatter_Matrix)
39
40     ## Eigen Vectors/ Values:
41     eig_val_sc, eig_vec_sc = np.linalg.eig(Scatter_Matrix)
42
43     ## Check if Su= u :
44     for i in range(len(eig_val_sc)):
45         eigv = eig_vec_sc[:,i].reshape(1,D).T
46         np.testing.assert_array_almost_equal(Scatter_Matrix.dot(eigv), eig_val_sc[i]*eigv,decimal=6,
47             err_msg='The eigenvector eigenvalue calculation is NOT correct.', verbose=True)
48
49     ## Rank the eigenvectors from highest to lowest corresponding eigenvalue and choose the top k
50         eigenvectors.
51
52     # Make a list of (eigenvalue, eigenvector) tuples
53     Pairs = [(np.abs(eig_val_sc[i]), eig_vec_sc[:,i]) for i in range(len(eig_val_sc))]
54
55     # Sort the (eigenvalue, eigenvector) tuples from high to low __ Using lambda function :)
56     Pairs.sort(key=lambda x: x[0], reverse=True)
57     print(len(Pairs))
58     # Checking that the list is correctly sorted
59     # for i in Pairs:
60         print(i[0])
61     #print(len(Pairs))
62
63     ## Construction of eigenvector matrix U.
64
65     q = input("Please enter S if you wish to use SVD to calculate array U or E for Eigendecomposition:
66         ")
67
68     if q == 'S':
69         # (i) SVD:
70         U,S,V = np.linalg.svd(X, full_matrices=False)
71     else:
72         # (ii) Eigendecomposition:
73         U = np.empty([D,M])
74         for i in range(M):
75             U[:,i] = Pairs[i][1].reshape(D,1)
76
77     print('Matrix U:\n', U)
78     print(U.shape)
79
80     ## Transforming the samples onto the new subspace with M- Dimension:
81
82     Projected_Data = U.T.dot(X)
83     print("shape {}".format(Projected_Data.shape))      ## Should be DxM

```

```

84     ## Plots
85
86     if len(y) == 1:
87         if M == 2:
88             Practical_Plots(Projected_Data)
89         else:
90             D_Plots(Projected_Data)
91             Practical_Plots(Projected_Data)
92     else:
93         if M == 1:
94             one_Plots(Projected_Data)
95         elif M == 2:
96             one_Plots(Projected_Data)
97             Theoretical_Plots(Projected_Data,y)
98         else:
99             one_Plots(Projected_Data)
100             Theoretical_Plots(Projected_Data,y)
101             T_D_Plots(Projected_Data,y)
102
103
104     ##### Probabilistic Principal Component Analysis #####
105
106     def PPCA(Arr,y,K):
107         print("\n#####__Probabilistic Principal Component Analysis__#####\n")
108         sigma_sq = int(input("Please enter 0 if s_square is zero: "))
109
110         ## Mean Vector:
111         D = Arr.shape[0]
112         N = Arr.shape[1]
113
114         mean_vector = np.empty((D,1))
115         for i in range(D):
116             mean_vector[i] = np.mean(Arr[i,:])
117
118         mean_d = np.repeat(mean_vector,N,axis=1)      # The mean vector with DxN shape
119         X = Arr- mean_d                               # X = x - mean(x)   X: Array of the Data   DxN
120
121         ## Initializing the parameters:
122
123         #sigma_sq = rand(0,1)
124         W_old = np.array(np.random.rand(D,K))         # W_old: DxK
125         #print("The matrix W is {}".format(W_old))
126         print("The shape of matrix W is {}".format(W_old.shape))
127         W_new = np.empty((D,K))
128
129         RMSE = mean_squared_error(W_old, W_new)**0.5   # Root Mean Squared Error (RMSE)
130         #print(RMSE)
131
132         ##### EM Algorithm:
133         Times = int(input("\nPlease assign the number of iterations:"))
134
135         # Limit case of s^2 -> 0
136         if sigma_sq == 0:
137             W_all = np.empty((D,K))
138             for i in range(Times):
139                 RMSEdiff = 1
140                 while RMSEdiff > 10**(-7):
141                     RMSEold = RMSE
142                     Omega = (linalg.inv((W_old.T).dot(W_old))).dot((W_old.T).dot(X))      # E step
143                     #print(Omega.shape)
144                     W_new = (X.dot(Omega.T)).dot(linalg.inv(Omega.dot(Omega.T)))          # M step
145                     #print(W_new.shape)
146                     RMSE = mean_squared_error(W_old, W_new)**0.5
147                     W_old = W_new
148
149                     RMSEdiff = abs(RMSE - RMSEold)
150
151                     print(RMSEdiff)
152
153                 W_all+= W_new
154
155         # s^2 != 0
156         else:
157             W_all = np.empty((D,K))
158             for i in range(Times):
159                 sigma_sq = random() #; print(sigma_sq)
160                 sigma_sq_new = random()
161                 dif_sigma=1
162
163                 RMSEdiff = 1
164                 while RMSEdiff > 10**(-7) or dif_sigma > 10**(-8):
165                     RMSEold = RMSE
166                     M = (W_old.T).dot(W_old) + sigma_sq*(np.eye(K))
167                     E_Zn = ((linalg.inv(M)).dot(W_old.T)).dot(X)
168                     E_Zn_ZnT = sigma_sq*linalg.inv(M) + E_Zn.dot(E_Zn.T)
169
170                     W_new = X.dot(E_Zn.T).dot(linalg.inv(E_Zn_ZnT)) #;print(W_new.shape)
171

```

```

172
173     for i in range(N):
174         Trace = np.trace(E_Zn_ZnT.dot((W_new).reshape(K,D)).dot(W_new))
175         a_01 = (np.linalg.norm(X[:,i].reshape(1,D)))**2
176         a_02 = ((E_Zn[:,i].reshape(1,K)).dot(W_new.T).dot(X[:,i].reshape(D,1)))
177         sigma_sq_new += np.sum(a_01 -2*a_02 +Trace)
178
179     sigma_sq_new = sigma_sq_new /(N*D)
180
181     RMSE = mean_squared_error(W_old, W_new)**0.5
182     dif_sigma = abs(sigma_sq - sigma_sq_new)
183
184     W_old = W_new
185     sigma_sq = sigma_sq_new
186
187     RMSEdiff = abs(RMSE - RMSEold)
188     #print(RMSE)
189     print(RMSEdiff)
190     print(dif_sigma)
191     print("\n")
192
193     W_all+= W_new
194
195 W_mean = W_all/(Times)
196
197 ### SVD and Projection
198
199 U,S,V = np.linalg.svd(W_mean, full_matrices=False)
200
201 Projected_Data = U.T.dot(X)
202 print(Projected_Data.shape)
203
204
205 if len(y) == 1:
206     if K == 2:
207         Practical_Plots(Projected_Data)
208     else:
209         D_Plots(Projected_Data)
210         Practical_Plots(Projected_Data)
211 else:
212     if K == 1:
213         one_Plots(Projected_Data)
214     elif K == 2:
215         one_Plots(Projected_Data)
216         Theoretical_Plots(Projected_Data,y)
217     else:
218         one_Plots(Projected_Data)
219         Theoretical_Plots(Projected_Data,y)
220         T_D_Plots(Projected_Data,y)
221
222 ##### Kernel Method #####
223
224 def KERNEL(Arr,y,M):
225
226     kernel = input("Construct the Kernel Matrix:\n Press G for Gaussian, P for Polynomial and T for
227                     Tanget: ")
228
229     D = Arr.shape[0]
230     N = Arr.shape[1]
231
232     ## Mean Vector:
233
234     mean_vector = np.empty([D,1])
235     for i in range(D):
236         mean_vector[i] = np.mean(Arr[i,:])
237     # print(mean_vector.shape)
238
239     mean_d = np.repeat (mean_vector,N,axis=1)      # The mean vector with DxN shape
240     X = Arr- mean_d                                # Normalize the Data matrix, X: DxN
241
242     N = X.shape[1]
243     K = np.empty((N,N)) #; print(K.shape)
244
245     ### Constructing the Kernels:
246
247     def Polynomial(Arr,p):
248         for i in range(N):
249             for j in range(N):
250                 K[i,j] = (1 + np.inner(Arr[:,i],Arr[:,j]))**p
251
252     def Tanget(Arr,delta):
253         for i in range(N):
254             for j in range(N):
255                 K[i,j] = np.tanh(np.inner(Arr[:,i],Arr[:,j]) + delta)
256
257     def Gaussian_Kernel(Arr,gama):
258         for i in range(N):

```

```

259         for j in range(N):
260             K[i,j] = np.exp(-gama*(((Arr[:,i] - Arr[:,j]).T).dot(Arr[:,i] - Arr[:,j]))**2))
261
262     ## Kernels:
263
264     if kernel == 'G':
265         gama = float(input("\nGive value for gama: ")) #gama = int(input("\nGive value for gama: "))
266         Gaussian_Kernel(X,gama)
267     elif kernel == 'P':
268         p = float(input("\nGive value for p: ")) #p = int(input("\nGive value for p: "))
269         Polynomial(X,p)
270     else:
271         delta = float(input("\nGive value for delta: ")) #delta = int(input("\nGive value for delta: "))
272         Tanget(X,delta)
273
274     One_N = np.empty((N,N))
275     for i in range(N):
276         One_N[i] = 1/N
277
278     ## Method:
279
280     K_bar = K-(One_N.dot(K)) - (K.dot(One_N)) + ((One_N.dot(K)).dot(One_N))
281
282     eig_values, eig_vectors = np.linalg.eig(K_bar)
283     Pairs = [(np.abs(eig_values[i]), eig_vectors[:,i]) for i in range(len(eig_values))]
284
285     # Sort the (eigenvalue, eigenvector) tuples from high to low __ Using lambda function :)
286     Pairs.sort(key=lambda x: x[0], reverse=True)
287     print("length {}".format(len(Pairs)))
288
289     U = np.empty([N,M])
290     for i in range(M):
291         U[:,i] = Pairs[i][1].reshape(D,1)
292
293     print('Matrix U:\n', U)
294     print(U.shape)
295
296     Projected_Data = U.T.dot(K_bar)
297     print(Projected_Data.shape)
298
299     if len(y) == 1:
300         if M == 2:
301             Practical_Plots(Projected_Data)
302         else:
303             D_Plots(Projected_Data)
304             Practical_Plots(Projected_Data)
305     else:
306         if M == 1:
307             one_Plots(Projected_Data)
308         elif M == 2:
309             one_Plots(Projected_Data)
310             Theoretical_Plots(Projected_Data,y)
311         else:
312             one_Plots(Projected_Data)
313             Theoretical_Plots(Projected_Data,y)
314             T_D_Plots(Projected_Data,y)
315
316
317     ##### Function for Subquestion (ii) of Theoretical Exercise #####
318
319     def Theoretical_II(Arr01,Arr02):
320         print(Arr02.shape)
321         def Eigenvalues(Arr):
322             ## Mean Vector:
323             D = Arr.shape[0]
324             N = Arr.shape[1]
325
326             mean_vector = np.empty([D,1])
327             for i in range(D):
328                 mean_vector[i] = np.mean(Arr[i,:])
329             mean_d = np.repeat (mean_vector,N,axis=1) # The mean vector with 3x80 shape DxN
330             X = Arr- mean_d # Normalize the join matrix X DxN
331
332             ## Scatter Matrix: is used to estimate the Covariance matrix of a multivariate normal distribution
333             Scatter_Matrix = np.empty([D,D])
334             for i in range(X.shape[1]):
335                 Scatter_Matrix += (X[:,i].reshape(D,1)).dot((X[:,i].reshape(D,1)).T)
336
337             ## Eigen Vectors/ Values of Scatter Matrix:
338             eig_values, eig_vectors = np.linalg.eig(Scatter_Matrix)
339
340             val = np.ndarray.tolist(eig_values)
341             sort_val = sorted(val,reverse =True)
342
343             return sort_val
344
345
346

```

```

347 sort_val = Eigenvalues(Arr01)
348 sort_val2= Eigenvalues(Arr02)
349 print("Eig 10:{}".format(sort_val))
350 print("Eig 5:{}".format(sort_val2))
351 maxx = max([max(sort_val),max(sort_val2)])
352 x_val = [i for i in range(1,len(sort_val)+1)]
353
354 fig, ax = plt.subplots()
355
356 ax.scatter(x_val,sort_val,marker='o', color='blue', alpha=0.5, label='N=10')
357 ax.scatter(x_val,sort_val2,marker='*',color='red',alpha=0.5,label='N=5')
358 ax.set_ylim([0,maxx+10])
359
360 plt.grid(True)
361 plt.legend( numpoints=1 ,loc='upper right')
362 plt.title('Eigenvalues_Plot')
363 plt.savefig("Eigenvalues_Plot.png")
364 plt.show()
365
366
367
368 ##### Plot Functions #####
369
370 ## Theoretical Part:
371
372 def Theoretical_Plots(Arr,y):
373
374     plt.scatter(Arr[0,y==0], Arr[1,y==0],color='red',marker='^',alpha=0.5,label='Circle_01')
375     plt.scatter(Arr[0,y==1], Arr[1,y==1],color='blue',marker='o',alpha=0.5,label='Circle_02')
376     plt.grid(True)
377     plt.xlabel('Pca_01')
378     plt.ylabel('Pca_02')
379     plt.legend(numpoints =1,loc='lower right')
380     plt.title('Projection')
381     plt.savefig("Theoretical_01.png")
382     plt.show()
383
384 def T_D_Plots(Arr,y):
385
386     fig = plt.figure(figsize=(8,8))
387     ax = fig.add_subplot(111, projection='3d')
388     plt.rcParams['legend.fontsize'] = 10
389
390     ax.scatter(Arr[0,y==0], Arr[1,y==0],Arr[2,y==0],color='red',marker='^',alpha=0.5,label='Circle_01')
391     ax.scatter(Arr[0,y==1], Arr[1,y==1],Arr[2,y==1],color='blue',marker='o',alpha=0.5,label='Circle_02')
392     ax.grid(True)
393
394     ax.set_xlabel('Pc_01')
395     ax.set_ylabel('Pc_02')
396     ax.set_zlabel('Pc_03')
397
398     plt.title('Projected in 3D')
399     ax.legend(numpoints=1,loc='lower right')
400     plt.savefig("3d Plot.png")
401     plt.show()
402     plt.close()
403
404 def one_Plots(Arr):
405     y = [2 for i in range(500)]
406     plt.plot(Arr[0,0:500],y,'o', markersize=7, color='blue', alpha=0.5, label='Circle_01')
407     plt.plot(Arr[0,500:1000],y,'^', markersize=7, color='red', alpha=0.5, label='Circle_02')
408     plt.grid(True)
409     plt.xlabel('Pc_01')
410     plt.ylabel('Pc_02')
411     #plt.ylim([-10,10])
412     plt.legend(numpoints=1,loc='lower right')
413     plt.title('Projection in 1d')
414     plt.savefig("1d Plot.png")
415     plt.show()
416
417
418 ## Practical:
419
420 def Practical_Plots(Arr):
421
422     plt.plot(Arr[0,0:3], Arr[1,0:3], 'o', markersize=7, color='blue', alpha=0.5, label='Baseline')
423     plt.plot(Arr[0,3:27], Arr[1,3:27], '^', markersize=7, color='red', alpha=0.5, label='Normal Diet')
424     plt.plot(Arr[0,27:51], Arr[1,27:51], '*', markersize=7,color='green',alpha=0.5,label='High-fat Diet')
425     plt.grid(True)
426     plt.xlabel('Pc_01')
427     plt.ylabel('Pc_02')
428
429     plt.legend(numpoints=1,loc='lower right')
430     plt.title('Mice Projection')
431     plt.savefig("Mice Projection.png")
432     plt.show()
433
434

```



```

435 def D_Plots(Arr):
436
437     fig = plt.figure(figsize=(8,8))
438     ax = fig.add_subplot(111, projection='3d')
439     plt.rcParams['legend.fontsize'] = 10
440
441     ax.plot(Arr[0,0:3],Arr[1,0:3],Arr[2,0:3], 'o', markersize=7, color='blue', alpha=0.5,
442             label='Baseline')
443     ax.plot(Arr[0,3:27],Arr[1,3:27],Arr[2,3:27], '^', markersize=7, color='red', alpha=0.5,
444             label='Normal Diet')
445     ax.plot(Arr[0,27:51],Arr[1,27:51],Arr[2,27:51], '*',
446             markersize=7,color='green',alpha=0.5,label='High-fat Diet')
447
448     ax.grid(True)
449     ax.set_xlabel("Pc_01")
450     ax.set_ylabel("Pc_02")
451     ax.set_zlabel("Pc_03")
452     plt.title('Projected Mice 3D')
453     ax.legend(numpoints=1,loc='lower right')
454     plt.savefig("Mice 3d Plot.png")
455     plt.show()
456     plt.close()

```

5 Suggested Bibliography

References

- [1] Christopher M. Bishop *Pattern Recognition and Machine Learning*
- [2] I. T. Jolliffe *Principal Component Analysis, Second Edition*
- [3] Sebastian Raschka *Python Machine Learning*
- [4] Lingbo Yu, Robert R. Snapp, Teresa Ruiz and Michael Radermacher *Probabilistic Principal Component Analysis with Expectation Maximization (PPCA-EM) Facilitates Volume Classification and Estimates the Missing Data* <https://www.ncbi.nlm.nih.gov/pmc/articles/PMC3353830/>
- [5] Richard Everson, Stephen Roberts *Inferring the eigenvalues of covariance matrices from limited, noisy data* <http://empslocal.ex.ac.uk/people/staff/reverson/uploads/Site/spectrum.pdf>
- [6] César Souza, *Kernel Functions for Machine Learning Applications* http://crsouza.com/2010/03/17/kernel-functions-for-machine-learning-applications/#kernel_choosing